

Deep Learning

Vazgen Mikayelyan

YSU, Krisp

September 27, 2019

1 Stochastic Gradient Descent

2 Introduction to Tensorflow

Stochastic Gradient Descent

Let L be a loss function that we know:

Stochastic Gradient Descent

Let L be a loss function that we know:

$$L(w) = \frac{1}{n} \sum_{i=1}^n (f_w(x_i) - y_i)^2,$$

$$L(w) = \frac{1}{n} \sum_{i=1}^n (-y_i \log f_w(x_i) - (1 - y_i) \log (1 - f_w(x_i))),$$

$$L(w) = \frac{1}{n} \sum_{i=1}^n \left(-y_i^T \log f_w(x_i) \right).$$

Stochastic Gradient Descent

Let L be a loss function that we know:

$$L(w) = \frac{1}{n} \sum_{i=1}^n (f_w(x_i) - y_i)^2,$$

$$L(w) = \frac{1}{n} \sum_{i=1}^n (-y_i \log f_w(x_i) - (1 - y_i) \log (1 - f_w(x_i))),$$

$$L(w) = \frac{1}{n} \sum_{i=1}^n \left(-y_i^T \log f_w(x_i) \right).$$

Do you see problems in finding minimum of these functions using GD?

Note that in each case we can represent the loss function by the following form:

$$L(w) = \frac{1}{n} \sum_{i=1}^n L_i(w).$$

Note that in each case we can represent the loss function by the following form:

$$L(w) = \frac{1}{n} \sum_{i=1}^n L_i(w).$$

SGD algorithm is the following:

- Choose an initial vector of parameters w and learning rate α .

Note that in each case we can represent the loss function by the following form:

$$L(w) = \frac{1}{n} \sum_{i=1}^n L_i(w).$$

SGD algorithm is the following:

- Choose an initial vector of parameters w and learning rate α .
- Repeat until an approximate minimum is obtained.

Note that in each case we can represent the loss function by the following form:

$$L(w) = \frac{1}{n} \sum_{i=1}^n L_i(w).$$

SGD algorithm is the following:

- Choose an initial vector of parameters w and learning rate α .
- Repeat until an approximate minimum is obtained.
 - Randomly shuffle examples in the training set.

Note that in each case we can represent the loss function by the following form:

$$L(w) = \frac{1}{n} \sum_{i=1}^n L_i(w).$$

SGD algorithm is the following:

- Choose an initial vector of parameters w and learning rate α .
- Repeat until an approximate minimum is obtained.
 - Randomly shuffle examples in the training set.
 - For $i = 1, 2, \dots, n$, do $w \leftarrow w - \alpha L_i(w)$.

Note that in each case we can represent the loss function by the following form:

$$L(w) = \frac{1}{n} \sum_{i=1}^n L_i(w).$$

SGD algorithm is the following:

- Choose an initial vector of parameters w and learning rate α .
- Repeat until an approximate minimum is obtained.
 - Randomly shuffle examples in the training set.
 - For $i = 1, 2, \dots, n$, do $w \leftarrow w - \alpha L_i(w)$.

Do you see problems in this optimization method?

Mini-Batch Gradient Descent

MBGD algorithm is the following:

- Choose an initial vector of parameters w , learning rate α and batch size B .

Mini-Batch Gradient Descent

MBGD algorithm is the following:

- Choose an initial vector of parameters w , learning rate α and batch size B .
- Repeat until an approximate minimum is obtained.

Mini-Batch Gradient Descent

MBGD algorithm is the following:

- Choose an initial vector of parameters w , learning rate α and batch size B .
- Repeat until an approximate minimum is obtained.
 - Randomly shuffle examples in the training set.

Mini-Batch Gradient Descent

MBGD algorithm is the following:

- Choose an initial vector of parameters w , learning rate α and batch size B .
- Repeat until an approximate minimum is obtained.
 - Randomly shuffle examples in the training set.
 - For $i = 1, 2, \dots, \lceil \frac{n}{B} \rceil$, do

$$w \leftarrow w - \alpha \frac{1}{B} \sum_{k=(i-1) \cdot B + 1}^{i \cdot B} L_k(w).$$

Outline

1 Stochastic Gradient Descent

2 Introduction to Tensorflow

Difference Between



CPU



GPU



TPU

CPU vs GPU vs TPU

CPU vs GPU vs TPU

- Central Processing Unit is the electronic circuitry, which work as a brain of the computer that perform the basic arithmetic, logical, control and input/output operations specified by the instructions of a computer program.

CPU vs GPU vs TPU

- Central Processing Unit is the electronic circuitry, which work as a brain of the computer that perform the basic arithmetic, logical, control and input/output operations specified by the instructions of a computer program.
- The Graphics Processing Unit is a specialized electronic circuit designed to render 2D and 3D graphics together with a CPU. GPU also known as Graphics Card in the Gammer's culture. Now GPU are being harnessed more broadly to accelerate computational workloads in areas such as financial modeling, cutting-edge scientific research, deep learning, analytics and oil and gas exploration etc.

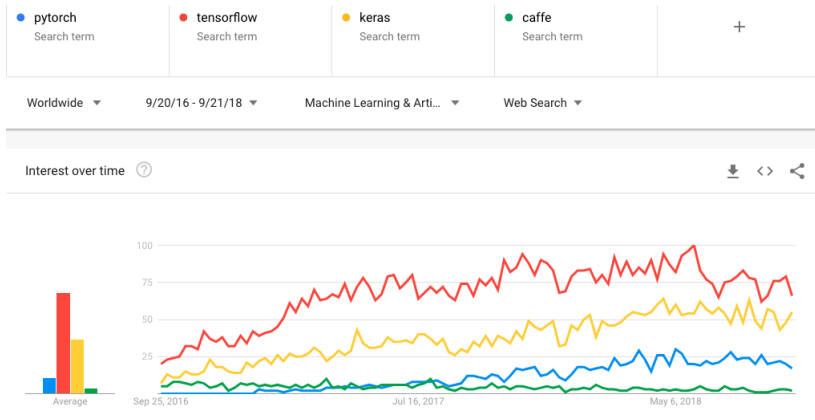
CPU vs GPU vs TPU

- Central Processing Unit is the electronic circuitry, which work as a brain of the computer that perform the basic arithmetic, logical, control and input/output operations specified by the instructions of a computer program.
- The Graphics Processing Unit is a specialized electronic circuit designed to render 2D and 3D graphics together with a CPU. GPU also known as Graphics Card in the Gammer's culture. Now GPU are being harnessed more broadly to accelerate computational workloads in areas such as financial modeling, cutting-edge scientific research, deep learning, analytics and oil and gas exploration etc.
- Tensor Processing Unit is a custom-built integrated circuit developed specifically for machine learning and tailored for TensorFlow, Google's open-source machine learning framework. TPU's have been powering Google data centers since 2015.


Why Tensorflow?



Interest Over Time







Google Search Activity




GitHub Activity

GitHub Activity



Online Job Listings



What can do Tensorflow?

- TensorFlow lets you deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device in a very simple way. This way the things can be done very fast.

What can do Tensorflow?

- TensorFlow lets you deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device in a very simple way. This way the things can be done very fast.
- TensorFlow lets you express your computation as a data flow graph.

What can do Tensorflow?

- TensorFlow lets you deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device in a very simple way. This way the things can be done very fast.
- TensorFlow lets you express your computation as a data flow graph.
- TensorFlow lets you visualize the graph using the in-built tensorboard. You can inspect and debug the graph very easily.

What can do Tensorflow?

- TensorFlow lets you deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device in a very simple way. This way the things can be done very fast.
- TensorFlow lets you express your computation as a data flow graph.
- TensorFlow lets you visualize the graph using the in-built tensorboard. You can inspect and debug the graph very easily.
- TensorFlow gives the best performance with an ability to iterate quickly, train models faster and run more experiments.

What can do Tensorflow?

- TensorFlow lets you deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device in a very simple way. This way the things can be done very fast.
- TensorFlow lets you express your computation as a data flow graph.
- TensorFlow lets you visualize the graph using the in-built tensorboard. You can inspect and debug the graph very easily.
- TensorFlow gives the best performance with an ability to iterate quickly, train models faster and run more experiments.
- TensorFlow runs on nearly everything: GPUs and CPUs—including mobile and embedded platforms—and tensor processing units (TPUs), which are specialized hardware to do the tensor math on.

Basic Code Structure

- View functions as computational graphs.


Basic Code Structure

- View functions as computational graphs.
- First build a computational graph, and then use a session to execute operations in the graph.

Basic Code Structure

- View functions as computational graphs.
- First build a computational graph, and then use a session to execute operations in the graph.
- This is the basic approach, there is also a dynamic approach implemented in the recently introduced eager mode.

Computational Graphs



Computational Graphs

- Nodes are operators (ops), variables and constants.

Computational Graphs

- Nodes are operators (ops), variables and constants.
- Edges are tensors
 - 0-d is a scalar
 - 1-d is a vector
 - 2-d is a matrix
 - Etc.

Computational Graphs

- Nodes are operators (ops), variables and constants.
- Edges are tensors
 - 0-d is a scalar
 - 1-d is a vector
 - 2-d is a matrix
 - Etc.
- Constants are fixed value tensors - not trainable.

Computational Graphs

- Nodes are operators (ops), variables and constants.
- Edges are tensors
 - 0-d is a scalar
 - 1-d is a vector
 - 2-d is a matrix
 - Etc.
- Constants are fixed value tensors - not trainable.
- Variables are tensors initialized in a session - trainable / not trainable.



Computational Graphs

- Nodes are operators (ops), variables and constants.
- Edges are tensors
 - 0-d is a scalar
 - 1-d is a vector
 - 2-d is a matrix
 - Etc.
- Constants are fixed value tensors - not trainable.
- Variables are tensors initialized in a session - trainable / not trainable.
- Placeholders are tensors of values that are unknown during the graph construction, but passed as input during a session.

Computational Graphs

- Nodes are operators (ops), variables and constants.
- Edges are tensors
 - 0-d is a scalar
 - 1-d is a vector
 - 2-d is a matrix
 - Etc.
- Constants are fixed value tensors - not trainable.
- Variables are tensors initialized in a session - trainable / not trainable.
- Placeholders are tensors of values that are unknown during the graph construction, but passed as input during a session.
- Ops are functions on tensors.

Computational Graphs



- Session is the runtime environment of a graph, where operations are executed and tensors are evaluated.

- Session is the runtime environment of a graph, where operations are executed and tensors are evaluated.

```
>>> import tensorflow as tf
>>> a = tf.constant(1)
>>> b = tf.constant(2)
>>> op1 = tf.add(a, b)
>>> print(a)
Tensor("Const:0", shape=(), dtype=int32)
>>> print(b)
Tensor("Const_1:0", shape=(), dtype=int32)
>>> print(op1)
Tensor("Add:0", shape=(), dtype=int32)
```

- Session is the runtime environment of a graph, where operations are executed and tensors are evaluated.

```
>>> import tensorflow as tf
>>> a = tf.constant(1)
>>> b = tf.constant(2)
>>> op1 = tf.add(a, b)
>>> print(a)
Tensor("Const:0", shape=(), dtype=int32)
>>> print(b)
Tensor("Const_1:0", shape=(), dtype=int32)
>>> print(op1)
Tensor("Add:0", shape=(), dtype=int32)
```

```
>>> import tensorflow as tf
>>> a = tf.constant(1)
>>> b = tf.constant(2)
>>> c = tf.add(a, b)
>>> with tf.Session() as sess:
...     print(sess.run(a))
...     print(sess.run(b))
...     print(sess.run(c))
...
1
2
3
```

- Session is the runtime environment of a graph, where operations are executed and tensors are evaluated.

```
>>> import tensorflow as tf
>>> a = tf.constant(1)
>>> b = tf.constant(2)
>>> op1 = tf.add(a, b)
>>> print(a)
Tensor("Const:0", shape=(), dtype=int32)
>>> print(b)
Tensor("Const_1:0", shape=(), dtype=int32)
>>> print(op1)
Tensor("Add:0", shape=(), dtype=int32)
```

```
>>> import tensorflow as tf
>>> a = tf.constant(1)
>>> b = tf.constant(2)
>>> c = tf.add(a, b)
>>> with tf.Session() as sess:
...     print(sess.run(a))
...     print(sess.run(b))
...     print(sess.run(c))
...
1
2
3
```

- `a.eval()` is equivalent to `session.run(a)`, but in general, “eval” is limited to executions of a single op and ops that returns a value.

- Session is the runtime environment of a graph, where operations are executed and tensors are evaluated.

```
>>> import tensorflow as tf
>>> a = tf.constant(1)
>>> b = tf.constant(2)
>>> op1 = tf.add(a, b)
>>> print(a)
Tensor("Const:0", shape=(), dtype=int32)
>>> print(b)
Tensor("Const_1:0", shape=(), dtype=int32)
>>> print(op1)
Tensor("Add:0", shape=(), dtype=int32)
```

```
>>> import tensorflow as tf
>>> a = tf.constant(1)
>>> b = tf.constant(2)
>>> c = tf.add(a, b)
>>> with tf.Session() as sess:
...     print(sess.run(a))
...     print(sess.run(b))
...     print(sess.run(c))
...
1
2
3
```

- `a.eval()` is equivalent to `session.run(a)`, but in general, “eval” is limited to executions of a single op and ops that returns a value.
- Upon op execution, only the subgraph (required for calculating its value) is evaluated

Structure of Training Code

1. Assembling the graph
 - Create placeholders.

Structure of Training Code

1. Assembling the graph

- Create placeholders.
- Create variables / neural network.

Structure of Training Code

1. Assembling the graph

- Create placeholders.
- Create variables / neural network.
- Define a loss function.

Structure of Training Code

1. Assembling the graph

- Create placeholders.
- Create variables / neural network.
- Define a loss function.
- Define an optimizer.

Structure of Training Code

1. Assembling the graph

- Create placeholders.
- Create variables / neural network.
- Define a loss function.
- Define an optimizer.

2. Training in a session

- Start a session.

Structure of Training Code

1. Assembling the graph

- Create placeholders.
- Create variables / neural network.
- Define a loss function.
- Define an optimizer.

2. Training in a session

- Start a session.
- Initialize variables / restore from checkpoint.

Structure of Training Code

1. Assembling the graph

- Create placeholders.
- Create variables / neural network.
- Define a loss function.
- Define an optimizer.

2. Training in a session

- Start a session.
- Initialize variables / restore from checkpoint.
- Run the optimizer over batches.

TensorBoard

